


A decorative border of Egyptian hieroglyphs surrounds the central text. The hieroglyphs are arranged in vertical columns on the left and right sides, and in horizontal rows at the top and bottom. The symbols include various birds (falcon, ibis, duck), eyes, lotus flowers, ankh symbols, and other traditional Egyptian motifs.

الأبجدية المصرية

— El Abgadeyya El Maşreyya —

MASRI

Write Egyptian. One consistent orthography.

Diaa Hassouna

2026

El Abgadeyya El Masreyya — The Masri Writing System

A Language-Standard Specification

Diaa Hassouna

Contents

Part I — What Masri Is - 5

Chapter 1: The Case for This Approach - 5

Chapter 2: Definition - 6

Chapter 3: Scope - 7

Chapter 4: Design Goals - 7

Chapter 5: Reference Dialect - 7

Chapter 6: Canonicalization - 8

Part II — The Canonical Orthography (Masri Standard · Tier 2) - 9

Chapter 7: The Tier 2 Alphabet - 9

Chapter 8: Phonology vs. Orthography - 10

Chapter 9: Core Spelling Rules - 11

Chapter 10: Canonical Letter-to-Sound Table - 12

Chapter 11: Relationship to the Other Representations - 12

Part III — Representations - 13

Chapter 12: One System, Six Representations - 13

Chapter 13: Masri Basic (Tier 1) - 13

Chapter 14: DRS — The Romanization - 14

Chapter 15: Arabic-P — Plain Arabic - 15

Chapter 16: Arabic-V — Vocalized Arabic - 16

Chapter 17: Arabic-S — Simplified Arabic Vowel Script - 17

Chapter 18: Choosing a Representation - 17

Part IV — Language Standard - 19

Chapter 19: Vowels - 19

Chapter 20: Numbers - 20

Chapter 21: Morphology - 21

Chapter 22: Grammar - 25

Chapter 23: Syntax - 26

Part V — Lexical Standardization - 27

Chapter 24: Native Vocabulary - 27

Chapter 25: Loanwords - 27

Chapter 26: Proper Names and Brands - 29

Chapter 27: Variants and Attestation - 29

Part VI — Input and Implementation - 30

Chapter 28: Franco Normalization - 30

Chapter 29: Arabic Normalization - 30

Chapter 30: Mixed Input - 32

Chapter 31: Converter Architecture - 32

Chapter 32: Keyboards - 32

Chapter 33: Data Structures - 32

Part VII — Historical Context - 34

Chapter 34: Historical Egyptian Scripts - 34

Chapter 35: Earlier Writing Proposals for Modern Egyptian - 34

Chapter 36: Franco - 35

Chapter 37: eg-lang - 35

Part VIII — Development History - 37

Chapter 38: Handwritten Drafts - 37

Chapter 39: Design Decisions - 37

Chapter 40: Revisions - 38

Chapter 41: Roadmap - 38

Part I — What Masri Is

Chapter 1: The Case for This Approach

Before any alphabet, any grammar rule, or any conversion table, a fair question deserves an answer: why does Modern Egyptian need a new writing system at all? Four gaps in everything that came before make the case.

1.1 Why Arabic Script, Not Latinization

Every serious prior attempt to give Modern Egyptian a consistent written form — from Wilhelm Spitta's 1880 grammar through Carl Vollers, John Selden Willmore, and the more recent wave of romanization projects up through eg-lang — has reached for the Latin alphabet. Franco, the informal Arabizi that Egyptians actually use daily on phones and keyboards, is Latin-based too, extended with digits standing in for letters Latin can't represent.

This is a strange inheritance. The overwhelming majority of Egyptians read and write Arabic script fluently; it is what they are taught from childhood, what appears on every sign and every schoolbook, what their hands already know how to form. A writing system that asks them to abandon that script in favor of a Latin one is asking for more than spelling reform — it is asking for a change of visual habit that most of the population has no practical reason to make. Latinization has never lacked linguistic sophistication; what it has lacked is a plausible route to everyday use. Masri starts from the opposite assumption: the script Egyptians already read is the script Egyptians should write in. Where a Latin-based representation is still useful — for typing on basic keyboards, for academic transliteration, for continuity with existing Franco habits — Masri provides one. But it is offered as one option among several, not as the foundation.

1.2 The Standardization Gap

Ask ten Egyptians to write the same colloquial sentence in Arabic script and you will likely get ten different spellings. Should a long vowel be marked or left to inference? Should a word borrowed from French keep its foreign spelling or be respelled to match how it actually sounds? Should the sentence be written the way it would appear in a Standard Arabic sentence with the same root, or the way it's actually pronounced in the street? There is no single answer in current practice, because there is no standard — only a loose, self-taught convention that varies by writer, by region, and by how much Modern Standard Arabic spelling habit a given writer carries over into colloquial writing. Franco has exactly the same problem in Latin script: the same word appears with different digit choices, different vowel lengths, different consonant clusters, depending entirely on who is typing.

A language spoken by tens of millions of people every day has no agreed way to put itself on a page. That gap, more than any aesthetic argument about scripts, is the actual problem this book sets out to solve.

1.3 Why Arabic Script Alone Isn't Enough

Choosing Arabic script solves the familiarity problem, but it does not by itself solve the standardization problem. Ordinary Arabic script, written the way Egyptians already write it, systematically omits the information that would make a colloquial spelling unambiguous. Short vowels go unmarked. Gemination is inconsistently shown. There is no established convention for representing sounds that exist in Modern Egyptian but not in Modern Standard Arabic. Writing Modern Egyptian consistently in Arabic script, in other words, requires real spelling reform — not a new script, but new and explicit rules for an old one.

There are two established ways to carry out that kind of reform, and Masri proposes both, side by side, rather than forcing a choice between them:

- **The Arabic methodology** — full vocalization with ʔarakāt, the diacritical marks Arabic script already has for exactly this purpose, applied consistently and without exception. This is **Arabic-V**.
- **The Sorani-Kurdish methodology** — giving vowels their own ordinary letters, so that pronunciation is carried by the base script itself rather than by marks added above and below it. This is **Arabic-S**.

Both stay inside Arabic script. Both are fully reversible against the same canonical source. Neither is presented as the one correct answer; they are two disciplined solutions to the same underlying problem, offered to two different kinds of reader.

1.4 The Missing Integrated Standard

Grammar and vocabulary are not new territory — dictionaries and grammars of Egyptian Arabic exist, some of them excellent, going back well over a century. What has not existed, in any of the prior efforts this book is aware of, is a single orthography built to work consistently in both directions of script at once: a system where a word's spelling in Arabic script and its spelling in Latin script are not two separate traditions loosely related by transliteration, but two faces of one canonical form, mechanically derivable from each other. Grammar told in isolation from a fixed spelling standard doesn't give a language a stable written form; it gives it a set of observations about a form that still doesn't exist. Masri is built the other way around: fix the canonical spelling first, then let every representation — Arabic or Latin, marked or unmarked — derive from it.

Chapter 2: Definition

Modern Egyptian (اللغة المصرية الحديثة) is the language spoken as a first language by the majority of Egyptians in daily life, with its own grammar, phonology, and a substantial layer of vocabulary distinct from Modern Standard Arabic.

(The full argument for Modern Egyptian's status as a language in its own right — including the historical throughline traced in Part VII, Chapter 34 — is developed there rather than asserted here. This chapter states the claim; Part VII supplies the evidence.)

Masri is the name of the writing standard this book defines for that language: a single canonical orthography, expressed through six interoperable representations, covering script, phonology, grammar, and lexicon.

Chapter 3: Scope

This standard covers:

- A canonical phonemic inventory and its orthographic encoding (Part II)
- Six representations of that encoding, for different scripts and different needs (Part III)
- Grammar, morphology, and numeral conventions (Part IV)
- Policy for native vocabulary, loanwords, and proper names (Part V)
- Input methods and conversion between Franco, Arabic, and Masri (Part VI)

It does not claim to legislate how any individual chooses to write informally, does not propose replacing Modern Standard Arabic in religious, legal, or pan-Arab contexts, and does not assert political independence for Modern Egyptian as a matter of policy — its argument is linguistic and practical, not political.

Chapter 4: Design Goals

- **Precision** — every representation must be capable of encoding pronunciation without ambiguity, even where a given representation chooses not to show every detail by default.
- **Reversibility** — conversion between representations must be mechanical and lossless back to the canonical form.
- **Script parity** — Arabic-script and Latin-script representations are treated as equally legitimate expressions of the same standard, not as an original and a translation.
- **Teachability** — rules should be statable simply enough to teach, not just precise enough to formalize.
- **Borrowing, not reviving** — where Masri draws on Coptic or Greek letterforms, it does so as a source of visually appropriate, historically resonant characters for a new purpose, not as a claim of linguistic continuity or revival.

Chapter 5: Reference Dialect

Masri's canonical forms are anchored to **Cairene** Modern Egyptian. A written standard requires a single reference pronunciation to canonicalize against; without one, “correct” spelling has no fixed target. Cairene is chosen for the same reason most national standards choose a capital-region variety: it is the most widely understood variety across the country and the one most speakers can already parse even when it isn't their own. Speakers of other regional varieties are not asked to change how they speak — only, where they choose to use Masri's canonical spelling, to write against the Cairene reference the way a Scots English speaker writing “standard” English does not respell every word to match Scots pronunciation.

Chapter 6: Canonicalization

Every word in Masri has exactly one canonical form, defined in **Masri Standard (Tier 2)** — the subject of Part II. The other five representations described in Part III are not independent spelling systems that happen to resemble each other; they are deterministic derivations of that single canonical form, each surfacing a different amount of phonetic detail through a different script. This is what allows a word to be written in Arabic-P for everyday typing, in Arabic-V for precise study, in DRS for a basic Latin keyboard, or in Arabic-S as a middle path — and still be recognizably, checkably, the same word underneath.

Part II — The Canonical Orthography (Masri Standard · Tier 2)

Chapter 7: The Tier 2 Alphabet

Masri Standard is the canonical form every other representation in Part III derives from. It is a Latin-extended alphabet, built the way Coptic itself was built: where Coptic took the Greek alphabet as its base and added a handful of Demotic letters for sounds Greek couldn't represent, Masri takes the Latin alphabet as its base and adds a handful of restored Coptic and Greek letters for sounds Latin can't represent. The result is a script that reads as unmistakably Egyptian rather than as a diacritic-marked variant of English or French transliteration.

The alphabet follows the traditional hiġa'i letter order of Modern Standard Arabic, with minimal modification: the four letters not present in MSA (P, V, J, C) are each inserted beside their nearest phonetic or graphemic relative rather than reordering the whole system — P after B, V after F, J after G (G is the Cairene hard ج; J is /ʒ/, always written ج), and C after S.

Complete Tier 2 letter inventory (37 letters)

#	Letter	Arabic	Name	IPA	Origin
1	A/a	ا	Alef	/ʔ/, /α/, /εə/	ASCII
2	B/b	ب	Beh	/b/	ASCII
3	P/p	پ	Pee	/p/	ASCII
4	T/t	ت	Teh	/t/	ASCII
5	Θ/θ	ث	Ṭeh	/θ/ → /s/	Greek
6	G/g	ج	Geem	/g/	ASCII
7	J/j	چ	Jéh	/ʒ/	Extra
8	Ƶ/ƶ	ح	ƶa	/ħ/	Coptic
9	X/x	خ	Ƙa	/x/	ASCII
10	D/d	د	Daal	/d/	ASCII
11	Ð/ð	ذ	Ḍaal	/ð/ → /z/	Latin ext.
12	R/r	ر	Reh	/r/	ASCII
13	Z/z	ز	Zeen	/z/	ASCII
14	S/s	س	Seen	/s/	ASCII
15	C/c	—	See	/s/, /k/	Extra
16	ʃ/ʃ	ش	Sheen	/ʃ/	Coptic
17	Ṣ/ṣ	ص	Ṣaad	/sˤ/	Latin ext.
18	Ḍ/ḍ	ض	Ḍaad	/dˤ/	Latin ext.
19	Ṭ/ṭ	ط	Ṭah	/tˤ/	Latin ext.
20	ƶ/Ʒ	ظ	ƶah	/ðˤ/	Latin ext.

#	Letter	Arabic	Name	IPA	Origin
21	Ⲁ/Ⲁ	ع	‘Ain	/ʕ/	Coptic
22	Ⲅ/Ⲅ	غ	Gheen	/ɣ/	Latin ext.
23	F/f	ف	Feh	/f/	ASCII
24	V/v	ڤ	Vee	/v/	ASCII
25	Q/q	ق	Qaaf	/q/	ASCII
26	K/k	ك	Kaaf	/k/	ASCII
27	L/l	ل	Laam	/l/	ASCII
28	M/m	م	Meem	/m/	ASCII
29	N/n	ن	Noon	/n/	ASCII
30	H/h	ه	Heh	/h/	ASCII
31	O/o	و	Oh	/u/, /u:/	Vowel
32	U/u	و	Ū	/ʌ/	Vowel
33	W/w	و	Dablew	/ʔw/, /w/	Vowel
34	E/e	ي	Ee	/e/, /e:/, /ɛ/, /ɛ:/	Vowel
35	I/i	ي	Ai	/i/	Vowel
36	Y/y	ي	Wye	/ʔj/, /j/	Vowel
37	Ⲁ/Ⲁ	ء	Hamza	/ʔ/	Latin ext.

Notes on the entries that need more than a table row:

- **G vs. J** — G is the Cairene hard /g/ realization of ج; J (/ʒ/) covers both Upper Egyptian (Ṣa‘idi) ج and the loanword /ʒ/, always written ج in Arabic script.
- **P and V** — full, independent letters, never interchangeable with B or F. P is used mainly in loanwords and modern vocabulary; V is an original, coexisting native phoneme, not a loanword-only import (Chapter 15 corrects an earlier overstatement of this point).
- **C** — has no Arabic-script counterpart; restricted to code-switched words.
- **Ⲁ (Old Coptic Ain)** — descends from the hieroglyph for “eye.”
- Plain ه is ordinary ASCII H/h; there is no separate Coptic letter for plain /h/. Only the emphatic ح uses the Coptic-derived letter (Ⲅ/Ⲅ).

Chapter 8: Phonology vs. Orthography

Masri deliberately keeps two questions separate that colloquial Egyptian writing habits usually blur together: *how a word is pronounced* and *how a word is spelled*. This distinction is what makes the rest of the standard possible.

Phonology is the underlying system of sounds and how they pattern — which sounds exist in Modern Egyptian, which are allophones of each other, which contrasts are meaningful. Orthography is the decision about how those sounds get written down. The two are related but

not identical: an orthography can choose to under-represent phonological detail (as Arabic-P does, by omitting most ẓarakāt) without that choice implying anything false about the phonology itself. Conversely, a phonologically precise transcription (like Arabic-V or DRS at full detail) is not a *better* description of the language than a phonologically sparse one — it is a different orthographic choice about how much of the same underlying phonology to show on the page.

Keeping this distinction explicit is what allows Masri to define one canonical phonological form per word and then let six different orthographic representations surface different amounts of that same information. Without the distinction, “correct spelling” and “correct pronunciation” collapse into the same question, and — as Chapter 1 argued — that collapse is exactly what has kept Modern Egyptian without a stable written form for so long.

Chapter 9: Core Spelling Rules

A working set of rules that apply across the canonical form, independent of which representation eventually surfaces it:

- **The definite article *el*** is always written as a separate word, space-separated, never merged or hyphenated to the word it modifies (*el naharda*, not *elnaharda*).
- **Word-initial glottal stop** — whether from hamza or from a colloquial ق — is never written, even when the word follows *el*. The glottal belongs to the word itself, not to the sentence, so القمر is *el amar*, not *el Vamar*.
- **Sentence-initial capitalization** mirrors English convention: the first letter of a sentence is capitalized.
- **No capital letter appears mid-word or word-finally**, under any circumstance — including for special letterforms like < and ʔ, which are capitalized only at true word-initial position.
- **Gemination** is shown by doubling the consonant letter in Tier 2, the same principle Arabic-P shows with a shadda and Arabic-V shows with a shadda plus explicit vowel marking.
- **A word’s attested pronunciation is canonical**, not its Modern Standard Arabic cognate’s pronunciation. Where colloquial merger has collapsed a Modern Standard Arabic distinction (as ث does, unpredictably, into either *t* or *s* depending on the word), the canonical Masri form follows the actual Egyptian sound, memorized per word rather than derived by rule.
- **Cairene ق and hamza are treated as one phoneme**. Both realize as the glottal stop ∇ in everyday Cairene speech, so both are spelled identically as ∇ where the glottal is written at all (mid-word/final, per the rule above) — *ba’ool* (everyday) / *ba∇ool* (Tier 2). *Q* is reserved separately, for formal/MSA-register pronunciation only (*Qur’an*, never *∇ur’an*).
- **P and V are full, independent letters**, never treated as interchangeable with B or F respectively, regardless of how a given loanword or native word’s nearest MSA cognate would be spelled.

Chapter 10: Canonical Letter-to-Sound Table

The letter-to-sound correspondence is the table already given in full in Chapter 7 — every Tier 2 letter, its Arabic-script counterpart, its IPA value, and its origin. It is not repeated here as a second table, since two copies of the same authoritative data is exactly the kind of drift this book has tried to avoid throughout (Chapters 20, 25, 35–37 all defer to a single source the same way). What this chapter adds instead is the vowel-specific detail Chapter 7’s consonant-oriented table doesn’t carry — see Chapter 19 for the complete vowel-to-symbol correspondence, including short/long pairing and the haraka-vs-carrier-letter distinction Arabic-V depends on.

Chapter 11: Relationship to the Other Representations

Every representation described in Part III is a deterministic function of the canonical Tier 2 form defined in this Part — never an independent spelling decision made from scratch. Concretely:

- **Masri Basic (Tier 1)** substitutes each non-ASCII Tier 2 letter with an ASCII-safe equivalent, following fixed rules (ع → backtick; the digraphs th/kh/dh/sh/gh for their respective Coptic-letter equivalents; emphatics as plain letter + caret).
- **DRS** romanizes the canonical form using keyboard-safe Latin-Extended substitutes for the handful of Coptic/Greek glyphs, everything else identical to Tier 2.
- **Arabic-P, Arabic-V, and Arabic-S** all use the same underlying Arabic-script letter inventory — including the extra consonants and vowel letters Tier 2 already defines — and differ only in how much of the canonical vowel and gemination detail each one chooses to mark explicitly.

Because every representation traces back to the same Tier 2 source, converting a word from any one representation to any other is never a matter of interpretation — it is a lookup against the canonical form.

Part III — Representations

Chapter 12: One System, Six Representations

Masri is not six competing writing systems. It is one canonical language standard — the Tier 2 form defined in Part II — surfaced through six different representations, each built for a different script, a different keyboard, and a different reader.

The six representations split into two groups — a distinction the live reference data makes explicit and this book adopts directly:

Canonical representations — complete enough to serve as a primary reading/writing form in their own right, not just a fallback:

Representation	Role	Script	Vocalization
Masri Standard (Tier 2)	Primary	Latin-extended	Full, by letter choice
DRS	Companion	Latin-Extended	Full, keyboard-safe substitutes
Arabic-V	Companion	Arabic script + ʔarakāt	Full
Arabic-S	Companion	Arabic script + 2 extra letters	Full, via vowel letters

Practical fallbacks — built for constrained input contexts, not for full-precision reading:

Representation	Script	Vocalization
Masri Basic (Tier 1)	ASCII Latin	Reduced
Arabic-P	Arabic script	Minimal (shadda only)

A note on this grouping: the live site's own representations data currently lists only five representations (Tier 2, DRS, Arabic-V, Tier 1, Arabic-P) — Arabic-S predates that data's last update and isn't yet registered there. When it is, "companion" alongside DRS and Arabic-V is the natural fit, matching the table above. [ROADMAP ITEM]

Every one of the six is a deterministic derivation of the Tier 2 canonical form, never an independent spelling decision — the mechanism Chapter 11 already established. The canonical/fallback split is about completeness of what a representation can express on its own, not about correctness — a fallback representation is exactly as accurate as a canonical one for what it chooses to show; it simply shows less.

Chapter 13: Masri Basic (Tier 1)

Masri Basic exists for one purpose: to be the direct, ASCII-safe alternative to Franco on a basic Latin keyboard, with no non-standard characters and no academic apparatus. Where Tier 2

reaches for a restored Coptic letter or a diacritic to represent a sound precisely, Tier 1 substitutes a plain-ASCII equivalent by fixed rule:

- ع → the backtick (`)
- ث خ ذ ش غ → the digraphs *th, kh, dh, sh, gh*
- Emphatic consonants → the plain letter plus a caret (*s[^], d[^], t[^], z^{and h}* for ح — plain *h* stays *h*, unmarked)
- A hyphen separates letters that would otherwise misread as one of these digraphs (*Mas-h[^]oor*, not *Mashoor*)
- The glottal stop is written ' only mid-word or word-final — never word-initial, consistent with the core spelling rule from Chapter 9
- No digit ever stands in for a letter — Franco's digit-for-letter habit (3 for 7 ,ع for ح) is treated purely as legacy input, never as valid Tier 1 output

Tier 1 is not a simplified version of the language — it carries exactly the same phonemic content as Tier 2, letter for letter, just without any character a basic keyboard can't produce. Anyone who has ever typed Franco can read Tier 1 with almost no adjustment; the difference is that Tier 1 spelling is fixed and predictable where Franco's never was.

Complete Tier 1 substitution table

Tier 2	IPA	Tier 1
◀	/ʕ/	`
Ⲫ	/ħ/	h [^]
ⲱ	/ʃ/	sh
Ⲑ / Ⲕ	/θ/ / ð/	th / dh
Emphatic ⲥ ⲧ ⲛ ⲟ	pharyngealized	s [^] d [^] t [^] z [^]
Ⲳ	/ʔ/	' (mid/final only, never initial)
P V J	/p/ /v/ /ʒ/	p v j (unchanged — already ASCII)

Rule notes: a hyphen breaks a false digraph reading (*Mas-h[^]oor*); no digit ever substitutes for a letter, even where Franco habit would suggest one.

Chapter 14: DRS — The Romanization

DRS stands for Dīaa Romanization System — named for its author, in the same tradition as Buckwalter transliteration (Chapter 35), which likewise carries its creator's name. DRS is a Latin-script representation built for readers already comfortable with informal Egyptian typing. Mechanically, DRS is **Tier 2 minus five Greek/Coptic-block letters**, each replaced by a Latin-Extended lookalike that's visually close to the letter it replaces. Everything else in the Tier 2 inventory — the plain Latin base, P/V/J, Ⲳ, Ⲑ/ð, the emphatics — carries into DRS unchanged, which is what keeps DRS automatically in sync with Tier 2 rather than drifting into its own parallel spelling tradition.

Complete DRS substitution table

Tier 2 letter	Sound	DRS substitute	Why
Θ / θ	/θ/	Ʀ / Ƨ	Greek letter not on standard keyboards
Ϣ / ϣ	/ħ/	Ĥ / ĥ	Coptic letter not on standard keyboards
Ϝ / ϝ	/ʃ/	Ŝ / ŝ	Coptic letter not on standard keyboards
ϥ (single-case)	/ʎ/	Ǝ (single-case)	Coptic letter not on standard keyboards; DRS keeps the same one-glyph-both-cases convention
Ϻ / ϻ	/ʒ/	Ğ / ğ	Greek letter not on standard keyboards

The Franco-like *feel* of DRS comes from those five substitutes reading as ordinary accented Latin rather than an unfamiliar Unicode block — not from DRS being an independently-designed romanization tradition of its own.

Chapter 15: Arabic-P — Plain Arabic

Arabic-P is how the overwhelming majority of Egyptians already read and write Modern Egyptian in Arabic script, with only the small adjustments the canonical standard requires for consistency.

What it is. Arabic-P is not a new script — it is ordinary Arabic script, written the way Egyptians already type it day to day, with no unfamiliar apparatus imposed on the reader.

2arakāt policy. Arabic-P uses shadda (ّ) when gemination needs to be shown, and nothing else. Sukūn is omitted entirely, exactly as it is in everyday Egyptian typing — a deliberate simplification, not an oversight.

Full letter inventory. Arabic-P mirrors Masri Standard (Tier 2) letter-for-letter, including the consonants and vowel letters not present in Modern Standard Arabic (پ, چ, ڦ; ئ, ۆ where the writer chooses to use them). One distinction worth being precise about: **V is an original, coexisting native Egyptian phoneme, not a loanword import** — پ is used mainly for loanwords and modern vocabulary, but ڦ represents a sound Modern Egyptian already has on its own. Neither is interchangeable with its nearest MSA neighbor (پ ≠ ب, ڦ ≠ ف), and a word using either is written with the letter that sound actually is, rather than being forced through an MSA-only inventory that has no clean way to represent it.

Because Arabic-P, Arabic-V, and Arabic-S all share the same full Tier 2-derived letter set, moving between any of the three Arabic-script representations is a matter of adding or removing marking, not relearning a script.

Chapter 16: Arabic-V — Vocalized Arabic

Arabic-V is the **Arabic methodology** answer to the problem posed in Chapter 1: full, exception-free ʔarakāt marking, applied with the same rigor Classical Arabic grammar applies to case-ending vocalization — but directed at representing Modern Egyptian precisely rather than at marking Modern Standard Arabic grammatical function.

Who it’s for. Arabic-V is built for readers already fluent in Arabic script who want precision — the ability to read a word once and know exactly how it’s pronounced, with nothing left to inference. It trades typing speed for that certainty.

The fundamental principle. Every vowel is marked, every gemination is marked with a shadda, and — where the standard’s vowel-length and vowel-quality distinctions require it — a small high seen (◌ِ, U+06DC) disambiguates a consonant-plus-fatᶇa sequence from an open-vowel digraph, so that ڤِ reads unambiguously as *ya* while unmarked ڤ reads as *é*, and ڤِ reads as *wa* while unmarked ڤ reads as *ou*.

Consonant treatment. Arabic-V shares Arabic-P’s full letter inventory and marks gemination explicitly and consistently.

A point of agreement, and a point of departure. Sami Harak’s own writing on Arabic script — quoted directly in the Grammar Reference — makes the case that Modern Egyptian has no grammatical need for vocalization at all, since it is analytic and carries no *i’rāb*, visible or estimated. Arabic-V agrees with that argument completely. It marks vowels anyway, not to represent grammatical case the way Classical Arabic ʔarakāt historically did, but as pure phonetic notation for a readership that wants full precision — the same principle Harak himself states, that every sound should be visible in writing, just applied through a different tool for a different purpose than his own plain-Arabic approach.

Arabic-V vowel reading table

Sequence	Reading	Note
ڤِ (yeh + fatᶇa + high seen)	<i>ya</i>	No exceptions
ڤ (yeh + fatᶇa, unmarked)	<i>é</i>	As in <i>café</i>
ڤِ (waw + fatᶇa + high seen)	<i>wa</i>	No exceptions, including <i>ywaddi</i>
ڤ (waw + fatᶇa, unmarked)	<i>ou</i>	Vowel of “toe”
Plain ڤ with ḍamma	<i>o</i>	As in “look”
◌ُ (ḍamma) standalone	<i>u</i>	Foreign words only (<i>menu</i> → منيو)
◌, word-final isolated	<i>-a</i>	—
◌ + suffix	<i>-et</i>	—
Sukūn on final l	non-geminated final <i>a</i>	—
Long vowels in use	<i>oo, ee, éé, aa, ou</i>	<i>ii</i> unused; <i>uu</i> is Canal-dialect, not Cairene

Chapter 17: Arabic-S — Simplified Arabic Vowel Script

Arabic-S is the **Sorani-Kurdish methodology** answer to the same problem: instead of marking vowels with *ṣarakāt*, give every vowel its own ordinary letter, so pronunciation is carried by the base script itself.

The vowel inventory.

Egyptian vowel	Arabic-S letter	Function
a / á	ا	Short/open a
e	ي	Front vowel, by position
i	ي	High front i
é / ε	ئ	Open-mid front vowel
o	و	Rounded o/u
ou / o:	ۆ	Mid-rounded, long o
glottal stop	ء	/ʔ/

Two dedicated letters. ئ (U+06CE, Arabic Letter Yeh With Small V) and ۆ (U+06C6, Arabic Letter Oe) are both already independent Arabic-script letters in Unicode — not combining marks — and both already see use in other languages' Arabic-based orthographies (Kurdish, Uyghur, Kazakh, and others for ۆ). Arabic-S repurposes them for Egyptian /ε/ and /o/ respectively, giving the system two vowels that need no diacritic at all: ئ as the front/open counterpart, ۆ as the rounded/back counterpart.

Core principle. Where traditional Arabic writes *consonant letters + optional ṣarakāt*, Arabic-S writes *consonant letters + independent vowel letters* — the vowel is never a hidden mark or an inferred pronunciation; it is an ordinary letter sitting in the word exactly where a consonant would sit.

Worked example. اسانسير (elevator) becomes ءاسانسير in Arabic-S, analyzable letter by letter as ر + ئ + ئ + س + ن + ا + س + ا + ء, with the doubled ي marking the long /ε:/.

Chapter 18: Choosing a Representation

No representation in this Part is more “correct” than another — each is a complete, faithful derivation of the same canonical Tier 2 form. The choice is entirely about audience and purpose:

- **Typing on a basic Latin keyboard, no special characters available** → Tier 1
- **Coming from Franco, want the closest thing to no learning curve** → DRS
- **Typing Arabic script the way you already do every day** → Arabic-P
- **Reading Arabic script and want zero ambiguity in pronunciation** → Arabic-V
- **Want explicit vowels in Arabic script without ṣarakāt** → Arabic-S
- **Writing by hand, or want the form that most feels like Egyptian's own script** → Tier 2

A single word can move through all six without ever leaving the standard — that portability, not any one representation's superiority, is the actual achievement Part II and Part III together are building toward.

Part IV — Language Standard

Chapter 19: Vowels

19.1 The short vowel inventory. Nine short-vowel entries, not five — the system is considerably richer than a simple five-way a/e/i/o/u split:

Tier 2	IPA	Example
a, á	[ɑ] open back, “father”	<i>Katab</i>
e	[e] front vowel	<i>We</i>
o	[u] close back rounded, “boot”	<i>borr</i>
ou	[o] close-mid back rounded, “go” (or [ɔ], “law”)	<i>loušt, Kours</i>
é	[ɛ] the Egyptian short é	<i>kaféhaat</i>
u	[ʌ] the fast “you” (yw) — foreign words only	<i>menu</i>
i	[i] word-initial or word-final only, never medial (a medial short front vowel is e)	<i>Soori, Roodi</i>
a (final)	final short a	<i>Huda</i>
o (final)	final short o	<i>kebro, kallemo, Leedo</i>

19.2 The long vowel inventory. Six entries:

Tier 2	IPA	Example
aa	[ɛə] the Egyptian long a	<i>Ketaab</i>
ee	[e:] long e	<i>sa <eed</i>
ii (final only)	[i:] long i — word-final only ; everywhere else, long i is written <i>ee</i>	<i>katabhomlii</i>
oo	[u:] long close back rounded	<i>boor</i>
ée	[ɛ:] the Egyptian long é	<i>asansээр</i>
uu	[o:] long close-mid, outside Cairo only (Canal dialects, Damietta) — not part of the Cairene reference standard (Chapter 5)	<i>buur</i>

19.3 Diphthongs. *ay* (a + y, y as a post-vocalic consonant: *Ray, Kay*) and *aw* (a + w: *law*).

19.4 Two long front vowels, one Arabic spelling. *ee* and *é/ée* are phonemically distinct vowels that Arabic script cannot tell apart on its own — both can surface as **ي** regardless of which is meant — but Masri spells them differently on purpose across every representation that can. This distinction is the only thing separating the ordinary sound-masculine plural suffix (*-een*, plain *ee*) from the professional-title dual suffix (Chapter 21) — *ostazéen* “two professors” versus a hypothetical *ostazeen* would be two different grammatical categories if Arabic script’s **ين** were the whole story.

19.5 Arabic-V’s realization rule. Every vowel has one Tier 2 letter and one Arabic-V form: a **ṣaraka** on the consonant writes a short vowel, and a carrier letter (**ي و ا**) with shadda and sukūn writes the corresponding long one; *w* and *y* immediately after a vowel are treated as consonants, not vowel letters. This is the general rule Chapter 16’s specific **ي/ي** and **و/و** readings are instances of.

19.6 Vowel shortening under suffixation. A stem’s long vowel shortens when certain suffixes attach, and stays intact before others — the split is conditioned by which suffix, not by any general rule about syllable count:

- **Shortens the stem:** *-ha*, *-na*, *-kom*, *-hom*, the sound-plural *-aat*, the dual *-een/-teen*, and the negative *-sh(i)*
- **Leaves the stem intact:** *-i*, *-ak/-ek*, *-o*

This one split governs three domains at once: *béét* “house” → *béthā* “her house” (nouns, Chapter 21); *Ostaaz* “professor” → *Ostazéen* “two professors” (duals, Chapter 21); *extaar* “he chose” → *ma xtarsh(i)* “he didn’t choose” (verbs under negation, Chapter 22).

19.7 Syncope. A short vowel in a word’s first syllable drops entirely when a suffix is added, if the resulting consonant cluster is pronounceable without it: *keber* “he grew up” → *ma kbersh(i)* “he didn’t grow up”; *de’ee* “flour” → *d’e’na* “our flour.” This is a deletion, not a spelling simplification — the vowel is genuinely absent from the spoken word, and Masri spells what’s said.

19.8 Epenthesis. Where deleting or shortening a vowel would leave three consonants in a row — a cluster Modern Egyptian doesn’t tolerate — an *e* is inserted to break it up instead: *kont* “I was” → *ma kontesh(i)* “I wasn’t,” not the unpronounceable *ma kontsh(i)*. Epenthesis and syncope are opposite repairs to the same underlying pressure; which one applies to a given word is lexically fixed, not predictable from the consonants alone.

Chapter 20: Numbers

Numerals in Modern Egyptian have their own native word-forms — *Talata*, not *Thalatha*; *Tamanya*, not *Thamaaneyā* — and, like the rest of colloquial Egyptian, have never had one standardized spelling before this reference. The rules below describe how the canonical Tier 2 forms are built; the complete 0-to-a-billion table itself is reproduced in full from the existing Numbers reference rather than retyped here.

Structure, by range:

- **0–10** are simple lexical items (*Ṣeḡr, WaṢed, Etneen, Talata...*), with no internal derivation.
- **11–19** are built by fusing the ones-digit onto *-(t)aṣar* (“-teen”): *Ṣedaṣar* (11), *Talattaṣar* (13), *Arba ṭaṣar* (14) — each ones-word contracted slightly before the suffix.
- **Tens (20, 30...90)** are their own derived forms ending in *-een* (◁*eṣreen* “20,” *Talateen* “30”), not compounds of the ones-digit.
- **21–99** compound as *[ones] we [tens]* — *WaṢed we ▷eṣreen* (“21”), literally “one and twenty,” mirroring the Arabic construction exactly.
- **Hundreds (200–900)** fuse the multiplier onto *-meyya* (“hundred”): *Toltomeyya* (300), *Xomsomeyya* (500) — 100 itself is simply *Meyya*, and 200 is irregular (*Meteeen*, a dual form, not “*Etneen meyya*”).
- **Thousands** work the same way against *-alaf*: *Talattalaf* (3,000), *Xamastalaf* (5,000); 1,000 and 2,000 are their own dual-marked forms (*Alf, Alfeen*).
- **Ten-thousands and hundred-thousands** compound as *[tens/hundreds] Alf* — ◁*eṣreen Alf* (20,000), *Meet Alf* (100,000).
- **Millions and billions** use borrowed/Arabic-derived standalone words: *Melyoon* (1,000,000), *Melyar* (1,000,000,000), combining with the smaller-number system the same way English does (“ten million” = ◁*aṣar Malayeen*).

Two rules carry through the whole system:

1. Every numeral is a *word*, spelled per Tier 2 rules like any other word — the Tier 1 rule from Chapter 13 that no digit ever stands in for a letter applies with equal force here.
2. Compounding follows the actual Cairene spoken pattern, not a written convention borrowed from Modern Standard Arabic.

Reference dialect: Urban Cairene, per Chapter 5.

Chapter 21: Morphology

21.1 Nouns. Modern Egyptian nouns carry no case marking and no nunation — the analytic profile already argued for in Chapter 16’s discussion of Harak’s *i’rāb* point carries through here as a morphological fact, not just a spelling convenience. A noun’s grammatical role is read entirely from word order and context, never from an ending. The one common survival is the frozen accusative in fixed expressions (شكراً *ṣokran* “thank you”).

21.2 Plurals. Egyptian plurals fall into two broad families — broken plurals (internal vowel-pattern change, the majority pattern for native vocabulary) and sound plurals (external suffix, *-een* for masculine, *-aat* for feminine and for most loanwords) — plus a small closed class of color/defect adjectives with their own plural template (*axdar/xadra/xodr* “green,” masc./fem./plural), and a genuine dual: masculine singular takes *-een*, the corresponding feminine takes *-teen* in place of its final *-a*. The dual is spelled with é (*Ostazéen*, not *Ostazeen*) to stay visually distinct from the ordinary sound-masculine plural, even though both look identical in Arabic script. Nisba adjectives of nationality/origin (*maṣri* “Egyptian”) skip the dual entirely and use the plural even for exactly two people — unlike profession nouns, which have a genuine dual.

21.3 Pronouns and clitics. Independent pronouns (*ana, enta, enti, howwa, heyya, eɓna, ento, homma*) are morphologically simple; the real complexity sits in the clitic system:

- **Possessive clitics on nouns** split into a “heavy” set (*-ha, -na, -kom, -hom*, and the sound-plural *-aat*) that shortens a preceding long stem vowel, and a “light” set (*-i, -ak/-ek, -o*) that leaves it intact — *béét* “house” → *béthā* “her house” but *bééti* “my house.”
- **Object clitics on verbs** stack in a fixed direct-then-indirect order (*ageeb* “I bring” → *agebhuulek* “I bring it to you”), with negation’s *-sh(i)* wrapping around the whole stack (Chapter 22).
- **Suffixed prepositions** (*fi-, <ala-/<a-, <and-*) conjugate like a defective verb paradigm. Shortened *<a-* fuses with a following *el* into *<al-*, the same pattern *li-* and *bi-* follow — but the comparison particle *ka* “like, as” deliberately never fuses with *el*, because a fused *kal* would be indistinguishable from *kal* “he ate.”

21.4 Verbs. Built from a 3- or 4-consonant root plus a vowel/affix pattern, cited in the 3rd-person-masculine-past form (Arabic has no infinitive) — *katab* “he wrote,” cited *katab/yekeṭeb*. Roots containing *ق/ك/ذ/ش/ع/ح/ت* take the same Tier 2 letter substitution in verb paradigms as anywhere else. **Hollow verbs** (middle radical *و/ي*, e.g. *gaab/yegeeb* “bring”) and **doubled verbs** (e.g. *ɓabb/yiɓebb* “love,” *baɗɗ/yeboɗɗ* “look”) each follow their own conjugation shape. The ten derived forms (I–X) remain productive; Form VII (*enbasat/yenbeset* “enjoy oneself”) and its VII_{It} variant (*etbasat/yetbeset*) are a documented case where two derivational shapes coexist as the same verb, and either spelling is correct.

Companion reference: verb paradigms by form

Every paradigm below applies the one part of verb inflection that’s fully attested and invariant across every form (Chapter 22’s *katab/yekeṭeb* affix template: *-t, -t, -ti, —, -et, -na, -tu, -u* for the past; *a-, te-, te-...-i, ye-, te-, ne-, te-...-u, ye-...-u* for the present subjunctive) to each form’s own attested citation stem. The affixes are mechanical; the stems are the real, sourced citation forms already given in §21.4/§4.6. Nothing below is invented — but the non-3rd-person cells are **derived**, not independently attested word-by-word, and should be checked against real usage before being treated as final.

Form I — strong root (*katab/yekeṭeb* “write,” already given in full at §22.1’s tense table)

Form I — with ɓ (*ɓabb/yiɓibb* “love”)

	Singular	Plural
Past 1st	ɓabbéét	ɓabbeena
Past 3rd m	ɓabb	ɓabbu
Pres. 3rd m	yiɓibb	yiɓibbu
Pres. indic. 3rd m	beyeɓebb	beyeɓebbu

Form I — with ɛ (*<eref/ye<raf* “know”)

	3rd m
Past	<eref

	3rd m
Present indic.	beye<raf
Participle	<aaref

Form I — hollow (*gaab/yegeeb* “bring,” middle radical ي/و)

	3rd m
Past 1st (I)	gebt
Past 3rd m	gaab
Present indic.	beygeeb

Form I — doubled (*başş/yeboşş* “look”)

	1st	3rd m
Past	başşet	başş
Present indic.	—	beyboşş

Form II (*darres/yedarres* “teach,” from §21.4/§4.6)

	Singular	Plural
Past 1st	darrest	darresna
Past 3rd m	darras	darrasu
Past 3rd f	darraset	—
Pres. subj. 1st	adarres	nedarres
Pres. subj. 3rd m	yedarres	yedarresu
Pres. indic. 3rd m	beyedarres	beyedarresu
Future 3rd m	hayedarres	—
Imperative	darres/darresi	darresu

Form III (*zaaker/yezaaker* “study with”)

	Singular	Plural
Past 1st	zaakert	zaakerna
Past 3rd m	zaaker	zaakeru
Pres. subj. 3rd m	yezaaker	yezaakeru
Pres. indic. 3rd m	beyezaaker	beyezaakeru

Form V (*etkallem/yetkallem* “speak,” reflexive of II)

	Singular	Plural
Past 1st	etkallemt	etkallemna
Past 3rd m	etkallem	etkallemu

	Singular	Plural
Pres. subj. 3rd m	yetkallem	yetkallemu
Pres. indic. 3rd m	beyetkallem	beyetkallemu
Negated (§22.2, stem-initial elision)	ma tkallemsh(i)	—

Form VII (*enbasaṭ/yeṇbeseṭ* “enjoy oneself,” passive of I; VIIIt variant *etbasaṭ/yeṭbeseṭ* equally correct, §21.4)

	Singular	Plural
Past 3rd m	enbasaṭ	enbasaṭu
Pres. subj. 3rd m	yeṇbeseṭ	yeṇbesṭu
Pres. indic. 3rd m	beyenbeseṭ	beyenbesṭu

Form VIII (*eštara/yešteri* “buy”)

	Singular	Plural
Past 3rd m	eštara	eštaru
Pres. subj. 3rd m	yešteri	yešteru
Pres. indic. 3rd m	beyešteri	beyešteru

Form X (*esta* <mal/*yesta* <mel “use”)

	Singular	Plural
Past 3rd m	esta<mal	esta<malu
Pres. subj. 3rd m	yesta<mel	yesta<melu
Pres. indic. 3rd m	beyesta<mel	beyesta<melu

Forms IV, VI, IX — not included. No attested Egyptian citation example for any of these three exists anywhere in this project’s sourced material. Form IV in particular is worth flagging rather than silently omitting: it’s largely non-productive in Modern Egyptian, having merged into Form I historically — a real linguistic fact, not a gap in research, though it should be confirmed against a dedicated source before this book states it as settled. Forms VI and IX simply have no attested example on hand yet; per Chapter 27’s evidentiary standard, that means they stay out rather than get a guessed citation word.

Root-type variation on derived forms (II, III, V, VII, VIII, X) is not yet covered above — every table there uses a strong root. How a hollow, doubled, defective, or assimilated root behaves under Form II’s or Form VIII’s template specifically (as opposed to Form I’s, which is fully attested above) is a real gap, and the honest thing to do is flag it rather than extrapolate from Form I’s behavior, since derived-form root-type interactions don’t always mirror Form I’s.

Chapter 22: Grammar

22.1 Tenses. Modern Egyptian marks tense/aspect through a small set of prefixes layered onto the present-subjunctive base:

- **Past** — the bare form (*katabt, katab...*).
- **Present subjunctive** — the bare base form (*akteb, tekteb, yekteb...*).
- **Present indicative** — subjunctive + *bi-*, always in its fuller *beye-* shape for 3rd person (*beyekteb*), never a reduced *b-* form.
- **Future** — subjunctive + *ha-* (*hakteb, hanekteb*).
- **Imperative** — its own short form (*ikteb/iktebi, iktebu*).

Compound tenses build on *kaan/yekoon* “was/will be” (an irregular, hollow-root paradigm in its own right): *kaan beyekteb* “he was writing,” *kaan katab* “he had written,” *haykoon katab* “he will have written,” *kaan hayekteb* “he was about to write.” The resultative/perfect uses the active participle standing in for a verb: *howwa kaateb* “he has written it.”

22.2 Negation. Modern Egyptian’s core negation pattern is *ma* (a free, separately-written word) plus *-sh(i)* (bound to the verb and its entire object-clitic stack) — not a circumfix wrapped as a single unit around the verb. *Ma* is written separately in every case, the same free-particle treatment Chapter 9 already established for *el*; only *-sh(i)* fuses:

- *ma katabsh(i)* “he didn’t write” — *ma* separate, *-sh(i)* glued to the bare verb
- *ma katabhomleesh* “he didn’t write them to me” — *-sh(i)* glues to the far side of the full object-clitic stack (Chapter 21), not to the bare verb

Future and interrogative negation use a different word entirely: *mesh* stands alone before a future verb (*mesh hayekteb* “he won’t write”) or before a bare verb in negated questions (*mesh katab?* “didn’t he write?”) — *mesh* is never glued to anything.

A presentative/topicalizing *ma* is a different word that happens to look identical. *Maho* (from *ma + howwa*) and the topicalizing *ma* in *Mal mohemma di...* fuse with a following *el* into *Mal*, exactly as *li-*, *bi-*, *fi-*, and shortened *◀a-* do (Chapter 21). **Negation *ma* never undergoes this fusion**, regardless of what follows it.

Bare *ma*, without *-sh(i)*, is a separate valid pattern — oath and emphatic-swearing register: *◀alayya el ne ◀ma ma ◀amalt ʔaaga!* “I swear I didn’t do a thing!” None of the phonological effects below apply here, since those are specifically triggered by adding *-sh(i)*.

Phonological effects of adding *-sh(i)* — shortening, syncope, and epenthesis — are governed by the rules already established in Chapter 19 (§19.6–19.8).

Compound tenses negate through *kaan* alone: *ma kansh(i) beyekteb* “he wasn’t writing,” not a double-negated form. The resultative/participle construction negates differently again, with *mesh* directly preceding the participle: *howwa mesh kaateb* “he hasn’t written it.”

22.3 Honorifics. Modern Egyptian marks social register and respect through a set of address terms — *baʔa* (informal, male peer), *hanem* (archaic/humorous address to a woman of high standing), *ostaaz* (teacher/expert, generic respectful address, with its own dual/plural per

Chapter 21), *me* < *allem* (working-class male address), *ra* < *Ve*es (boss, skilled laborer's title), *ba* < *mohandes* (engineer/generic middle-class male address) — distinct from MSA's honorific conventions and documented here as its own short system, since the choices involved are sociolinguistic, not purely grammatical agreement.

Chapter 23: Syntax

23.1 Word order. Modern Egyptian is SVO in its default, unmarked order, unlike Classical/MSA's VSO — a fixed typological fact that does real work: because nouns carry no case marking (Chapter 21), word order is the *only* signal of grammatical role. This is the syntactic half of the argument Chapter 16 already made about vocalization: a language with no *i'rāb* needs a reliable word-order convention precisely because it has nothing else to fall back on.

23.2 Questions. Content questions (*wh*-questions) are formed **in situ** by default — the question word stays in the position the answer would occupy — though fronting is also grammatical, giving two valid orders for the same question (*raaʔ maʃr emta?* / *Emta raaʔ maʃr?* “when did he go to Egypt?”). Wh-in-situ is flagged as a likely feature inherited from the Coptic substratum — a documented Coptic-influence claim, not merely an assertion. Compare to MSA, which always fronts.

23.3 Agreement. The dual, morphologically marked on nouns (Chapter 21), **collapses into ordinary plural agreement** on the verb and on any modifying adjective — a dual noun does not trigger a separate dual verb form the way Classical Arabic's dual agreement system does (*el ostazéen el soreyyeen dool beyem* < *lu lel gam* < *a* “these two Syrian professors are walking,” plural verb agreement despite the dual noun).

Part V — Lexical Standardization

Chapter 24: Native Vocabulary

Native Modern Egyptian vocabulary needs no separate lexical policy — it is already fully governed by the standards this book has built. A native word's canonical spelling comes from Part II's Tier 2 rules; its inflected forms come from Part IV's morphology (Chapter 21); where a native word survives in more than one attested regional form, Chapter 27's variant-recording model applies to it exactly as it applies to loanwords.

Personal names of Egyptians themselves are native vocabulary, not proper-name loanwords. An Egyptian's own name is spelled the same way any other native word is spelled — per Chapter 21's morphology, with no dual-representation apparatus and no source-language reference form. Chapter 26's proper-names policy exists for a different problem (foreign names entering Egyptian speech); it does not apply to Egyptians naming themselves.

Chapter 25: Loanwords

Loanwords are where lexical standardization actually earns its keep, because this is the one area where a writer's instinct — spell it the way the source language spells it — actively fights the rest of the standard.

25.1 Design goals. Five commitments: encode the reference Cairene speaker's actual pronunciation, never the donor language's spelling; support both retained-foreign and Egyptianized pronunciations as equally legitimate outcomes, decided by evidence; give every loanword sense exactly one canonical spelling, with genuinely attested secondary pronunciations recorded as explicit variants; reuse the extra letters Tier 2 already has (P, V, J) rather than inventing new ones; model nativization as graded, not binary.

25.2 Core principle. Loanwords are held to the same standard as every other Masri word. When principles conflict, pronunciation outranks everything else. The standard's general preference for orthography over phonetics (*el* always spelled *el* regardless of assimilation, Chapter 9) governs a specific, closed set of native morphological spellings — it is not license to spell loanwords by their donor-language letters.

25.3 The seven-category status model. Every loanword is diagnosed into one of seven categories — fully integrated, common code-switched, technical/international, proper noun, brand/trademark, recent borrowing, and a seventh residual class — each carrying its own default spelling strategy, diagnosed via §25.9.

25.4 Retained vs. Egyptianized pronunciation. Governed entirely by attested pronunciation and lexical status, never by how English or French a word looks. Retained/foreign-leaning is the expected default for technical and naming contexts and very recent borrowings; Egyptianized is the expected default once a word's dominant attested pronunciation shows adapted phonology. The system records what's actually said — it never prescribes a target pronunciation.

25.5 Adaptation rules. Where Tier 2 already has a letter for a foreign phoneme, that contrast is preserved by default — /p/ stays *p*, not *b*; /v/ stays *v*, not *f* — and nativizes to the native consonant only once the word has also taken native derivational morphology (§25.6): *بيتزا* → *bitza* (citation, retains *p*) versus its derived feminine *bitzaya* (nativized to *b*) — citation form and derived form are separate lexical entries and may genuinely diverge.

Where no Tier 2 letter exists at all, three specific sub-rules apply, and none of them invent a new letter:

- **Foreign θ/ð → t/d**, not the native Cairene merger targets *s/z* — a loanword *th*-sound and the native ث/ذ letters are two different substitutions that must not be conflated (*ترمس* → *tarmas*, not *tarsas*).
- **/tʃ/ and /ŋ/ get no dedicated letter.** /tʃ/ is written as *t+ح* where retained, simplifying to plain *ح* once fully nativized (*sandawitsh* retains it; *حبيبي* “chips” has nativized away from it). /ŋ/ is plain *n* before a following velar, the same allophonic treatment native words already get. /ʒ/ is the one exception with its own letter — *ج*, shared with native *Ṣaʿidi* ج — though an already-nativized older French loan may instead have hard-*g*’d historically (*جراج* → *garaaj*).
- **Non-native consonant clusters are retained as pronounced**, not broken up with an inserted vowel the way native Egyptian phonotactics would (*بلاستيك* → *blastik*); an epenthesized variant is recorded only when it’s itself independently attested, never invented by rule. Stress is not marked orthographically for loanwords, same as for native words.

25.6 Morphological integration. There is no separate “loanword morphology” — once a borrowed noun or verb reaches full integration, the same rules already fixed in Chapter 21 apply directly. A word’s citation form and its derived forms are canonicalized independently and may genuinely diverge — a cited dish name may keep a retained *p* while its derived, fully nativized form takes *b*.

25.7 Decision algorithm. Five steps: (1) establish the reference speaker’s actual attested pronunciation; (2) check whether it already differs from the source language’s; (3) diagnose lexical status via the seven categories; (4) check for native derivational morphology, which can override a category assignment; (5) assemble the canonical spelling, recording a second genuinely common pronunciation as an explicit accepted variant.

25.8 English and other source languages. English is the dominant current donor language and gets its own worked treatment for that reason alone. The identical, language-agnostic policy applies without modification to loans from French, Turkish, Italian, or any other donor.

25.9 Nativization threshold. Modeled as graded across six axes — frequency, age of borrowing, phonological adaptation, grammatical integration, speaker perception, and pronunciation variability — each scored low/medium/high from attested usage. No single axis is decisive alone.

25.10 Conflict-resolution hierarchy. Five tiers, in order: attested actual pronunciation (never overridden); phonemic contrast Tier 2 already encodes; lexical identity and morphological

integration together; source-language recoverability (a tiebreaker only); visual resemblance to source spelling (lowest, essentially never decisive).

A note on scope: the full English-loanword mapping table, the other-source-languages table, the worked-example set, and the 30-word stress test are extensive reference material reproduced wholesale from the existing Loanword Policy document rather than retyped here. [ROADMAP ITEM]

Chapter 26: Proper Names and Brands

Foreign proper names and brand names follow a **structured dual representation**, not a free-for-all:

- The **canonical Tier 2 spelling** is the name’s actual local Egyptian pronunciation, attested — Pepsi → *bebsi*.
- A **source-language reference form** may appear alongside on first mention, purely for recoverability — documentation practice, not a competing spelling.
- **Exception:** a name whose entire social function is international recognizability — a company’s own registered logotype, a person’s own preferred spelling — may be cited in its original form when directly quoting or labeling that specific entity, the same way English prose still capitalizes the P in “iPhone” mid-sentence.

As Chapter 24 already establishes, this policy governs *foreign* names entering Egyptian speech; Egyptians’ own personal names fall outside this chapter entirely.

Chapter 27: Variants and Attestation

One canonical, attested-default spelling per word sense, plus explicitly flagged accepted variants where a second pronunciation is genuinely common — never silent, unlimited variation. Three alternative approaches were considered and rejected: a single canonical spelling with no variants falsely marks real attested variation as error; fully free variation is exactly the uncontrolled-spelling problem this Part exists to prevent; treating every word as having two default-canonical forms overcorrects, manufacturing variation for words that don’t actually have two live pronunciations.

A variant earns formal recording, not silent tolerance, by the same evidentiary standard as everything else in this book: it has to be attested, not assumed. (Chapter 22’s Form VII/VIIIt note and the *itkatab/inkatab* passive variant are both instances of this same principle, not separate rules.)

Applied concretely to loanwords: *بيتزا* → *bitza* (canonical) / *pitza* (accepted variant) — *bizza* is simply wrong, not an unlisted third option. *باسبور* → *passport* / *passpour* (canonical + accepted variant) — *bassport* and *basboor* are wrong. The line between “accepted variant” and “wrong” is exactly this evidentiary standard: attested common pronunciations get recorded; unattested hypothetical respellings don’t, no matter how plausible they might look on paper.

Part VI — Input and Implementation

Chapter 28: Franco Normalization

Franco normalization is the one input path that is fully implemented today. Because Franco was never one system (Chapter 1), normalization here means resolving many attested spellings of the same word down to one canonical form, not merely transliterating.

28.1 Digit reading.

Franco digit	Arabic	Tier 2	Tier 1
2	ق / ء (Cairene)	∇ (silent word-initial)	' (mid/final only)
3	ع	◀	`
3'	غ	⧵	gh
4	ش	ω	sh
5	خ	x	kh
6	ط	ṭ	t^
6'	ظ	ẓ	z^
7	ح	ʕ	h^
8 (alt.)	غ	⧵	gh
9	ص	ṣ	s^
9'	ض	ḍ	d^
' (mid/final)	ء	∇	,

28.2 Tier 1 digraph resolution. $kh \rightarrow x$, $sh \rightarrow \omega$, $gh \rightarrow \gamma$, $th \rightarrow \theta$ (or s if merger applies), $dh \rightarrow \delta$ (or z under the same merger), $s^{\wedge} \rightarrow \mathfrak{s}$, $d^{\wedge} \rightarrow \mathfrak{d}$, $t^{\wedge} \rightarrow \mathfrak{t}$, $z^{\wedge} \rightarrow \mathfrak{z}$, $h^{\wedge} \rightarrow \mathfrak{h}$.

28.3 The hyphen rule. Where a digraph letter meets h or $h^{\wedge}$ and both consonants are genuinely pronounced, a hyphen keeps the reading unambiguous: مشهور $\rightarrow$ *Mashhoor* ($sh + h$), مسحور $\rightarrow$ *Mas-h^oor* ($s + h^{\wedge}$). Applies only at Tier 1; Tier 2's single-character letters make the ambiguity structurally impossible.

28.4 Variant collapsing. Every attested Franco spelling of a word maps to one canonical Tier 2 form: *3ashan*, *ashan*, *2ashan* $\rightarrow$ *◀aωan*. This is Chapter 27's variant policy in action.

28.5 Scope and limits. The converter is rule-based, not a language model. Genuinely ambiguous cases may need manual review against the standardized word list.

Chapter 29: Arabic Normalization

Arabic-script normalization is **not yet implemented in the live converter** — typing Arabic script in today passes the text through unchanged. What follows is a real design for it, built by inverting the deterministic Arabic-V generation rules already established (Chapter 16, the AV-R

rule set), rather than leaving this chapter as an unspecified aspiration. The three Arabic-script representations split cleanly into two cases, by how much information they actually carry.

29.1 Arabic-V input — fully deterministic, lossless. Because Arabic-V's generation from Tier 2 follows fixed rules (Chapter 16), every one of those rules can be read backwards to fully recover the canonical form, with no dictionary or guesswork needed:

- A ʔaraka directly on a consonant → the corresponding **short** vowel (Chapter 19's realization rule, inverted).
- A carrier letter (ا و ي) bearing shadda + sukūn → the corresponding **long** vowel.
- َ (yeh + fatḥa + high seen) → *ya*; unmarked ِ → *é*. ُ → *wa*; unmarked ُو → *ou*. (Chapter 16's reading table, run in reverse.)
- A shadda on any consonant → gemination, written as a doubled letter in Tier 2 (Chapter 9).
- Sukūn on a final ل → a non-geminated final *a*, per Chapter 19's table.

Because Arabic-V was designed from the start to make every phonemic detail explicit, this direction is a **pure lookup-table inversion** — the same category of operation as DRS↔Tier 2 (Chapter 14), not the harder problem Franco normalization solves.

29.2 Arabic-S input — also deterministic, also lossless. Simpler still, since Arabic-S carries its vowel information in ordinary letters rather than diacritics: each Arabic-S vowel letter (ا ي و) maps directly to one Tier 2 vowel per the table in Chapter 17, with no haraka/carrier distinction to resolve first. Doubled َؤ/َؤؤ reads as the corresponding long vowel exactly as doubled Tier 2 letters do.

29.3 Arabic-P input — the genuinely hard case. Arabic-P is minimally marked by design (Chapter 15) — only shadda survives, sukūn and short vowels are omitted entirely, exactly like ordinary Franco input. The consonant skeleton converts deterministically (a direct Tier 2 letter-substitution table, no ambiguity); the **vowels do not**, for the same structural reason Franco's digit-based input is ambiguous (Chapter 28): the information needed to reconstruct them simply isn't present in the input. This case requires the identical solution Chapter 28 already uses for Franco — a dictionary lookup against the standardized word list, falling back to the most statistically common vowel pattern for an unrecognized consonant skeleton, with the same manual-review caveat Chapter 28's §28.5 already states for genuinely ambiguous cases.

29.4 Why this hasn't shipped yet. §29.1 and §29.2 are close to a direct port of logic the Arabic-V generator already implements in reverse. §29.3 is the real remaining engineering work, and it's the same class of problem Franco normalization already solves — it needs the same dictionary infrastructure Chapter 28 relies on, applied to Arabic-script consonant skeletons instead of Franco digit-strings. Building this out remains the single highest-priority item on the project's roadmap (Chapter 41).

Chapter 30: Mixed Input

Mixed Franco+English text is handled today, normalized the same way pure Franco is. Arabic script mixed into that same input inherits Chapter 29's current limitation — it passes through unchanged rather than normalizing.

Chapter 31: Converter Architecture

Egyptian input (Arabic script / Franco / mixed)

↓
Normalization — rule-based, word-by-word

↓
Canonical Masri (Tier 2)

↓
Masri Standard · Tier 1 · DRS · Arabic-V · Arabic-P · Arabic-S

Three terms, used precisely:

- **Normalization / canonicalization** — resolving multiple attested input spellings to one preferred canonical form.
- **Romanization / representation** — re-expressing the already-canonical Tier 2 form in a different script for a different environment. The underlying word never changes, only the glyphs.
- **Transliteration** — mapping between writing systems while preserving the *source's* own spelling, the method scholarly Arabic transcription uses (Part VII). Masri's Arabic-to-canonical step is deliberately closer to normalization than transliteration.

Arabic-V's own generation follows this same pipeline: generated forward from canonical Tier 2, never an independent spelling-normalization system.

Chapter 32: Keyboards

- **Universal virtual keypad** — on-screen, no install, works on any device.
- **Windows keyboard layout** — installable driver, every extra letter on the key closest to its Arabic-keyboard counterpart, across four shift states.
- **Unicode hex entry** — type the code point, Alt+X (Word/Outlook).
- **Compose-key input** — WinCompose (Windows) or built-in OS compose sequences (macOS).
- **Text-replacement shortcuts** — iOS/macOS Text Replacement, Android/Gboard personal dictionary.
- **AnySoftKeyboard** — a dedicated Masri layer on Android.

Chapter 33: Data Structures

The system is data-driven: **alphabet.json** (the full Tier 2 letter inventory), **tier2-rules.json** (formal, numbered spelling rules), **drs.json** (the DRS substitution table, Chapter 14), and **loanword_examples** (the worked and stress-tested loanword entries backing Chapter 25) are structured reference files the converter, the alphabet reference, and this book itself all draw from.

Part VII — Historical Context

(This Part resolves the language-status flag from Part I, Chapter 2 — Chapter 34's timeline is the evidence for Modern Egyptian's status as a language in its own right.)

Chapter 34: Historical Egyptian Scripts

- **Old Kingdom (c. 3200–2100 BCE)** — no split: Old Egyptian spoken, governed, and written in together.
- **Middle Kingdom (c. 2100–1700 BCE)** — the first prestige/vernacular split: Middle Egyptian freezes into a literary standard that outlives its native speakers — the oldest ancestor of today's MSA/Modern-Egyptian gap.
- **New Kingdom (c. 1550–1070 BCE)** — the vernacular gets written down: Late Egyptian appears in everyday administration.
- **Late Period–Ptolemaic (c. 700–30 BCE)** — Demotic becomes the everyday hand; Ptolemaic rule brings Greek in for administration — ruling language and popular script diverge for the first time.
- **Roman Egypt (30 BCE–4th c. CE)** — Coptic script is born: spoken Egyptian starts being written with the Greek alphabet plus Demotic-derived signs for sounds Greek lacks. **The exact move Masri makes today.**
- **Byzantine/Coptic period (4th–7th c. CE)** — a full, identity-driven script change: Coptic becomes Christian Egypt's everyday script.
- **Early Islamic Egypt (641–9th c. CE)** — Arabic becomes the sole administrative language by the early 700s; Coptic stays spoken for generations after.
- **Medieval Egypt (10th–17th c.)** — Modern Egyptian emerges as the vernacular under Classical Arabic; Coptic dies out as everyday speech. From here, Modern Egyptian never holds official status.
- **Ottoman Egypt (1517–1867)** — a third layer (Ottoman Turkish, Classical Arabic, Modern Egyptian) — yet Arabic script remains the common thread throughout.
- **Khedival & British era (19th–20th c.)** — Latin script enters Egyptian daily life at scale for the first time, via French and English prestige, decades before Franco.
- **Contemporary Egypt (today)** — Modern Egyptian remains unofficial and spoken, written informally as Franco and in Arabic script; Masri proposes a standardized script for the vernacular itself.

The throughline: Modern Egyptian has never held official status at any point on this timeline. Egypt has changed its script for non-linguistic reasons — conquest, religion, identity — several times before. Masri isn't a break from that pattern; it's the next entry in it.

Chapter 35: Earlier Writing Proposals for Modern Egyptian

- **The Orientalist grammars (1880–1901)** — Spitta, Vollers, Willmore each invented a Latin transliteration for European students; Fiske's 1904 popularization of Spitta's system was elegant and phonetically airtight but needed diacritic type-sorts no Cairo print shop stocked. Masri shares Fiske's ambition but arrived at it through Franco, not

scholarship, and reaches into Egypt's own Coptic alphabet rather than bare Latin diacritics.

- **The al-Muqṭataf models (1889–1897)** — two rival Latin-script systems solving printing costs and mangled loanword spellings; neither stuck (an inverted letter “bothers the eye”). Masri borrows whole letterforms already in Egypt’s own alphabet instead.
- **The State Latinization Proposal (1944)** — Fahmi’s full Latin-replacement proposal, modeled on Atatürk’s 1928 reform, rejected within a month. His mistake was diplomatic, not linguistic — asking Egypt to trade its script for someone else’s. Masri’s Coptic letters are Egypt’s own, one chapter earlier.
- **Said Akl’s Lebanese alphabet (1961)** — proved a Latin script can carry serious dialect literature, but tied it to a Phoenician-descent political claim most readers didn’t share. Masri keeps the engineering, drops the politics.
- **The library romanization standards (1950s–1982)** — ALA-LC and DIN 31635, precise cataloguing systems for MSA/Classical Arabic with no attempt at spoken Cairene sounds. Beautiful precision, wrong job.
- **Buckwalter transliteration (1988)** — a strict, reversible ASCII mapping for Arabic NLP, still used in research. Brilliant for a computer, unreadable for a person.

Chapter 36: Franco

Franco/Arabizi has no single inventor — it emerged organically in the 1990s once pagers, early mobiles, and chatrooms supported only Latin characters, with numerals pressed into service for their loose visual echo of missing Arabic letters (3 for 7 ,ع for ح). It spread with zero institutional backing. Franco is proof of concept, not a finished product: it proved millions of Egyptians would rather invent a writing system than not write the dialect at all. **Masri's own lineage runs through this entry, not through the Orientalist grammars.**

Chapter 37: eg-lang

eg-lang, a contemporary project, proposes a fully Latin, diacritic-based orthography for Modern Egyptian — one alphabet, precisely tuned to Cairene phonology. The difference from Masri is architectural, not a matter of accuracy: eg-lang combines canonical orthography and representation into one single Latin system; Masri treats those as separate layers, with DRS serving the same practical need eg-lang addresses, offered alongside Tier 2 rather than replacing a dedicated script.

Sample comparison — eg-lang’s own opening line, and the same line in Masri (full homepage text belongs to eg-lang; read it at eg-lang.com/en):

eg-lang: “Ahlan biiku fi mašruugh ketaabet elgharabi elmašri belhuruuf ellatiniyya!”

Masri (Tier 2): Ahlan beeko fi maᵛrooᵼ ketaabet el ᵼarabi el Maṣri bel ʒoroof el Latineyya!

Even in one line, the architectural difference from §37's verdict is visible on the page: eg-lang leans on a dense stack of diacritics (ğ, š, š) to carry precision entirely within its one Latin alphabet, while Masri's Tier 2 line carries the same precision using dedicated letters (◀, ω, ʔ) rather than diacritic marks — the same “borrow whole letterforms rather than pile on

diacritics” principle Chapter 35 already credits Masri with over al-Muqtataf’s inverted-letter approach.

The scoreboard. Five axes, each attempt scored yes/partial/no against them — typable on an ordinary keyboard, naturally readable, captures Cairene phonology, carries a distinct Egyptian identity, actually adopted at scale:

Attempt	Typable	Readable	Captures Cairene	Egyptian identity	Adopted
Orientalist grammars	No	Partial	Partial	Partial	No
al-Muqtataf models	Partial	No	No	No	No
Fahmi Latinization	Yes	Partial	No	No	No
Said Akl’s alphabet	No	Yes	Partial	Partial	No
Library standards	No	No	No	No	No
Buckwalter	Yes	No	No	No	Partial
Franco	Yes	Partial	Yes	Partial	Yes
eg-lang	Partial	Yes	Yes	Partial	Partial
Masri	Yes	Yes	Yes	Yes	Yes

No prior attempt clears more than three of the five columns. Franco is the only one that was genuinely adopted, and it pays for that with weak readability and no real claim to Egyptian identity beyond convenience. Masri’s claim isn’t that any individual column is a novel achievement — Fahmi was typable, Said Akl was readable, Franco was adopted — it’s that no prior attempt cleared all five at once.

Part VIII — Development History

Chapter 38: Handwritten Drafts

Before there was a font, a keyboard layout, or a website, there was one handwritten page headed “Masri — Egyptian Alphabet.” It already lays out the rules the whole system still follows: one phonetic sound per letter, Latin as the base script filled out with Greek and Coptic letters where Latin runs short, the target being spoken Modern Egyptian rather than Classical Arabic — and, on principle, no digit is ever allowed to stand in for an Arabic letter, a deliberate break from Franco’s own numeral conventions from the very first page.

The idea came together around Eid al-Adha, 27 July 2026: a second handwritten page, signed and dated, proved the system by using it — writing the same sentence twice, once in Masri and once in Arabic underneath:

Masri: Hadaani tafkeeri waṽt <eed el Aḍḡa li waḍ< neḡaam ḡoroof <aṽaan kataabet el laḡa el Maṡreyya el Qaaḡereyya el <ammeyya el ḡaḍareyya bi estexdaam xaleeṡ men el ḡoroof el Lateneyya bel iḍaafa ila ṽewayyet ḡoroof Qebṡeyya we Yonaaneyya.

Arabic: هداني تفكيري وقت عيد الأضحى لوضع نظام حروف عشان كتابة اللهجة المصرية القاهرية العامية الحضرية باستخدام خليط من الحروف اللاتينية بالإضافة إلى شوية حروف قبطية ويونانية.

The page closes with a pointed correction, handwritten before there was any book to correct: the letter 2 (Coptic hori, standing in for ح) is not the digit 2 — a clear, deliberate difference the whole Tier 1 digit-reading system (Chapter 28) still enforces today.

Those two pages turned into the full 37-letter alphabet, the Tier 2 ruleset, and the public framework this book documents.

Chapter 39: Design Decisions

The founding principles, as stated in the project’s own reference data:

- **Regular sound-to-symbol correspondence.** Masri aims for a highly regular mapping between speech sounds and written symbols, with explicit orthographic rules — never silent exceptions — governing the contextual, morphological, and lexical cases where sound and symbol don’t map one-to-one (the θ→/s/ and ð→/z/ Cairene mergers, the ء/ق glottal merger, gemination).
- **Orthography over phonetics where the two conflict.** Some spellings preserve lexical or morphological identity rather than surface pronunciation — *el* is always written *el* regardless of sun/moon-letter assimilation in speech (Chapter 9). That is a deliberate orthographic choice, not a claim about actual pronunciation.
- **Simplicity over aesthetics** — speed and ease prioritized over formal elegance.
- **Cairene reference dialect** — urban Cairo Modern Egyptian as the normative standard (Chapter 5).

- **No numerals** — no Arabic or Western numeral ever substitutes for a letter, on principle, from the very first handwritten page (Chapters 13, 20, 28).
- **Keyboard accessibility at Tier 1** — every Tier 1 symbol typable on a standard Latin keyboard, with Tier 2 alone requiring supplementary Greek/Coptic input.
- **Hija’i letter order, minimally modified** — the alphabet follows MSA’s traditional letter order; letters absent from MSA are inserted beside their nearest phonetic relative rather than reordering the whole system (Chapter 7).
- **Reuse existing letterforms; invent nothing new.** Every non-Latin character in Tier 2 already existed somewhere before Masri repurposed it — the same rule behind Chapter 25’s loanword letters (P, V, J reused, never invented for loanwords specifically).
- **Reject architectures that overcorrect as readily as ones that undercorrect** — Chapter 27’s rejection of both “no variants” and “fully free variation,” and Chapter 25’s rejected dual-canonical-by-default loanword architecture.

Chapter 40: Revisions

- **The ε letter.** Earlier material referenced Greek Omega (Ω) for ε. The current letter is **Old Coptic Ain** (◄), adopted after a dedicated search of the Greek/Coptic Unicode block for the most suitable available letter (Chapters 7, 16, 17).
- **Negation spacing.** Earlier drafting treated *ma-...-sh(i)* as a single circumfix. The corrected rule (Chapter 22) writes *ma* as its own free word, exactly like *el*, while only *-sh(i)* stays bound to the verb and its clitic stack.
- **The *li-/bi-/fi-/◄a-/topicalizing-ma* fusion rule.** Formalized once confirmed across enough examples: this closed set fuses with a following *el*, while ◄*and* and negation *ma* never do.
- **Noun and verb vowel-shortening, unified.** A verb-specific observation was generalized into one rule covering nouns, duals, and the sound plural as well (Chapter 19, §19.6).

Chapter 41: Roadmap

Most of the source-pull items flagged earlier in this draft were resolved once the live reference data (`alphabet.json`, `tier2-rules.json`, `numbers.json`, `prior-attempts.json`) was checked against it — the complete 37-letter table (Chapter 7), the full vowel system (Chapter 19), the scoreboard (Chapter 37), and the detailed loanword rules (Chapter 25) are now sourced directly rather than flagged as pending. What’s genuinely still open:

- **Arabic-script normalization** (Chapter 29) — highest priority; the one piece of the architecture diagram that doesn’t yet do what the diagram says.
- **The verb-paradigm companion** (Chapter 21) now has a real first draft — Form I across all four root types, plus derived paradigms for Forms II, III, V, VII, VIII, and X — built by applying the attested, invariant person/tense template to each form’s sourced citation stem. What’s still open within it: Forms IV, VI, and IX have no attested Egyptian citation example anywhere in this project’s data, and root-type variation (hollow/doubled/defective/assimilated) on the derived forms hasn’t been verified beyond Form I. Both should be checked against real usage, not extrapolated, before being marked complete.

- **The eg-lang sample-line comparison** (Chapter 37) — eg-lang’s own opening line rendered in both systems; needs a source for eg-lang’s actual text, which wasn’t part of this project’s own reference data.
- **Registering Arabic-S in representations.json** (Chapter 12) — the live site’s own representations data doesn’t yet list Arabic-S alongside the other five.
- **A stray error in development-story.json** — the “Scrap of Paper” milestone’s own text currently reads “Greek omega for a long o,” left over from an early, since-corrected draft. The actual, current design decision — Old Coptic Ain (ϣ) for Ⲅ, not Omega for anything — is what Chapters 7, 16, 17, and 40 all correctly state; the site’s development-story data just hasn’t been updated to match. Worth a small source-data fix so the history page doesn’t contradict the alphabet page.